


# HPX

## A C++ Library for Parallelism and Concurrency

Hartmut Kaiser ([hkaiser@cct.lsu.edu](mailto:hkaiser@cct.lsu.edu))

CppCon, September 12, 2022

# A Real World Story





# Conventions

- Namespaces
  - `std::` namespace `std`
  - `stdex::` namespace `std::execution`
  - `stdexp::` namespace `std::experimental`
  - `hpx::` namespace `hpx`
  - `hpxex::` namespace `hpx::execution`
  - `hpxexp::` namespace `hpx::execution::experimental`
  - `hpxtt::` namespace `hpx::this_thread::experimental`

# HPX

The C++ Standards Library for Concurrency and Parallelism

<https://github.com/STELLAR-GROUP/hpx>

# HPX – An Asynchronous Many-task Runtime System

- At it's heart, HPX is a very efficient threading implementation
- Several functional layers are implemented on top:
  - C++ standards-conforming API exposing everything related to parallelism and concurrency
  - Full set of C++17/C++20/C++23 (parallel) algorithms
    - One of the first full openly available implementations
    - Extensions:
      - asynchronous execution
      - parallel range based algorithms
      - vectorizing execution policies `simd/par_simd`
  - Full set of senders/receivers (currently being discussed for standardization)
    - Implemented using C++17
  - Distributed operation
    - Extending the standard interfaces for use on tightly coupled clusters (super-computers)
    - Global address space, load balancing, uniform API for local and remote operations

# HPX – An Asynchronous Many-task Runtime System

- Full set of C++17/C++20/C++23 parallel algorithms

|                            |                          |                    |                         |
|----------------------------|--------------------------|--------------------|-------------------------|
| <u>adjacent_difference</u> | adjacent_find            | all_of             | any_of                  |
| copy                       | copy_if                  | copy_n             | count                   |
| count_if                   | equal                    | exclusive_scan     | fill                    |
| fill_n                     | find                     | find_end           | find_first_of           |
| find_if                    | find_if_not              | for_each           | for_each_n              |
| generate                   | generate_n               | includes           | inclusive_scan          |
| <u>inner_product</u>       | inplace_merge            | is_heap            | is_heap_until           |
| is_partitioned             | is_sorted                | is_sorted_until    | lexicographical_compare |
| max_element                | merge                    | min_element        | minmax_element          |
| mismatch                   | move                     | none_of            | nth_element             |
| partial_sort               | partial_sort_copy        | partition          | partition_copy          |
| reduce                     | remove                   | remove_copy        | remove_copy_if          |
| remove_if                  | replace                  | replace_copy       | replace_copy_if         |
| replace_if                 | reverse                  | reverse_copy       | rotate                  |
| rotate_copy                | search                   | search_n           | set_difference          |
| set_intersection           | set_symmetric_difference | set_union          | sort                    |
| stable_partition           | stable_sort              | swap_ranges        | transform               |
| uninitialized_copy         | uninitialized_copy_n     | uninitialized_fill | uninitialized_fill_n    |
| unique                     | unique_copy              |                    |                         |



# Parallel Algorithms

- Simple iterative algorithms
  - One pass over the input sequence
  - `for_each`, `copy`, `fill`, `generate`, `reverse`, etc.
- Iterative algorithms ‘with a twist’
  - One pass over the input sequence
  - Parallel execution requires additional operation after first pass, most of the time this is a reduction step
  - `min_element`, `all_of`, `find`, `count`, `equal`, etc.
- Scan based algorithms
  - At least three algorithmic steps
  - `inclusive_scan`, `exclusive_scan`, etc.
- Auxiliary algorithms
  - Sorting, heap operations, set operations, `rotate`

# Parallel Algorithms

- How does parallelization work?

- On CPUs

- Split input sequence into pieces (chunks) of theoretically arbitrary size
- Run algorithm on more than one core, each core on its own chunk
- Perform necessary synchronization and reduction

- On GPUs

- Split input sequence into pieces (chunks) that are sized to fit into a warp
- Run algorithm on more than one warp, each warp on its own chunk, each core on its own element
- Perform necessary synchronization and reduction

# Parallelize Loops

# Parallelize Loops

Sequence of elements:

|   |   |   |   |     |     |   |
|---|---|---|---|-----|-----|---|
| 0 | 1 | 2 | 3 | ... | N-1 | N |
|---|---|---|---|-----|-----|---|

```
std::vector<int> d = {...};  
hpx::for_each(d.begin(), d.end(), [](int val) {...});
```

---

```
template <typename Iterator, typename F>  
void for_each(Iterator b, Iterator e, F f)  
{  
    while (b != e)  
        f(*b++);  
}
```

# Execution Policies

- Standard introduces: `std::seq`, `std::par`, `std::unseq` (C++20), `std::par_unseq`
  - Passed as additional first argument to algorithm
- Convey guarantees/requirements imposed by loop body
  - `seq`: execute in-order (sequenced) on current thread
  - `unseq`: allow out-of-order execution (unsequenced) on current thread - vectorization
  - `par`: allow parallel execution on different threads
  - `par_unseq`: allow parallel out-of-order (vectorized) execution on different threads
- Proposed for standardization (P0350: Integrating `simd` with parallel algorithms): `stdex::simd`
  - Enable *explicit* vectorization that relies on special C++ types representing vector registers (`stdexp::simd`, see: Parallelism TS V2, latest draft: N4808)
- HPX introduces:
  - Asynchronous policies, e.g. `par(task)`: allow asynchronous operation
  - Explicit parallelized vectorization: `par_simd`
  - Executors: attached to execution policies using `.on()`

# Parallelize Loops

Sequence of elements:

|        |   |        |   |     |        |   |
|--------|---|--------|---|-----|--------|---|
| 0      | 1 | 2      | 3 | ... | N-1    | N |
| Core 0 |   | Core 1 |   |     | Core M |   |

```
std::vector<int> d = {...};  
hpx::for_each(par, d.begin(), d.end(), [](int val) {...});
```

---

```
template <typename Iterator, typename F>  
void for_each(parallel_policy, Iterator b, Iterator e, F f)  
{  
    auto size = std::distance(b, e);  
    std::vector<hpx::future<void>> v;  
    for (size_t chunk = 0; chunk != NUM_CHUNKS; ++chunk) {  
        v.push_back(hpx::async([&]() {  
            auto begin = std::next(b, (chunk * size) / NUM_CHUNKS);  
            hpx::for_each(begin, std::next(begin, size / NUM_CHUNKS), f);  
        }));  
    }  
    hpx::wait_all(v);  
}
```

// Iterator should be random access  
// assume: cleanly divisible  
// async() launches new thread, returns future  
// sequential for\_each()

# Parallelize Loops: Observations

- Parallelization concurrently runs sequential operations on parts of the input
  - At least for CPU based implementations
  - GPU based algorithms are usually different
- Iterators should be random access
  - Otherwise performance might be bad
- NUM\_CHUNKS is a magic number!
  - How should we select it?
  - What are the criteria for best performance?
- NUM\_CORES is another magic number
- AFFINITIES are important too (NUMA awareness!), control task placement

# A Bit of Background

Why is it so difficult to efficiently parallelize execution?



# Amdahl's Law (Strong Scaling)

$$S = \frac{1}{(1 - P) + \frac{P}{N}}$$

- S: Speedup
- P: Proportion of parallel code
- N: Number of processors






Figure courtesy of Wikipedia ([http://en.wikipedia.org/wiki/Amdahl's\\_law](http://en.wikipedia.org/wiki/Amdahl's_law))

# The 4 Horsemen of the Apocalypse: **SLOW**

- **S**tarvation
  - Insufficient concurrent work to maintain high utilization of resources
- **L**atencies
  - Time-distance delay of remote resource access and services
- **O**verheads
  - Work for management of parallel actions and resources on critical path which are not necessary in sequential variant
- **W**aiting for Contention resolution
  - Delays due to lack of availability of oversubscribed shared resources









# A Real World Story

## hpx::for\_each, Speedup, Intel Icelake, 64 cores





# Executors



# Executors

- Need abstraction of **How**, **When**, and **Where** to execute tasks
  - Need an API abstracting execution:
    - Fire & forget
    - Synchronously/asynchronously
    - Single/multiple tasks
    - Dependency tracking
  - Control parameters of execution
    - Chunk sizes?
    - Number of chunks?
    - What cores to use (where, number)?
    - Thread attributes (annotations, priorities, affinities, etc.)?

# Executors

- Executors abstract different task launching infrastructures
  - Synchronization using futures
    - HPX historically uses futures as main means of coordinating
  - Synchronization using sender/receivers (C++26?)
    - C++ standardization focusses on developing an infrastructure for anything related to asynchrony and parallelism
      - P2300: `std::execution` (senders & receivers)
    - Computational basis for asynchronous programming
    - Current discussions focus on integrating parallel algorithms
- In HPX, all facilities that launch tasks are implemented on top of executors
  - Parallel algorithms (attached to execution policies: `par.on(exec)`)
  - `hpx::async`, `hpx::dataflow`, `hpx::future::then`, etc. (passed directly as additional argument: `hpx::async(exec, f, ...)`)

# Executors: Parallel Algorithms

- HPX supports associating an executor with execution policies:

```
// Parallel execution using default executor
std::vector v = {1.0, 2.0, ... };
hpx::for_each(par, v.begin(), v.end(), [](double val) { ... });

// Parallel execution using parallel_executor
hpxex::parallel_executor exec;
hpx::for_each(par.on(exec), v.begin(), v.end(), [](double val) { ... });

// Parallel asynchronous (eager) execution using parallel_executor
future auto f = hpx::for_each(par(task).on(exec), v.begin(), v.end(), [](double val) { ... });
f.get(); // wait for completion

// Parallel execution using sender_executor
hpxexp::sender_executor sr_exec;
hpx::for_each(par.on(sr_exec), v.begin(), v.end(), [](double val) { ... });

// Parallel asynchronous (lazy) execution using sender_executor
sender auto s = hpx::for_each(par(task).on(sr_exec), v.begin(), v.end(), [](double val) { ... });
hpxtt::sync_wait(s); // start execution and wait for completion
```

# Executors: Parallel Algorithms

- HPX integrates parallel algorithms with senders/receivers

```
auto exec = ex::sender_executor();  
auto result =  
    hpxexp::just(std::begin(c), std::end(c), [](auto) { ... })  
    | hpx::for_each(par(task).on(exec))  
    | hpextt::sync_wait();
```

- Nicely integrates with existing Standard, does not require learning new APIs

# Executors

- HPX executors are (small) objects that expose an API supporting launching tasks:
  - `post` : fire & forget execution of given function
  - `sync_execute` : synchronously execute given function
  - `async_execute` : asynchronously execute given function, return awaitable
  - `bulk_async_execute` : asynchronously execute given function N times, return awaitable
  - `bulk_sync_execute` : asynchronously execute given function N times
  - `then_execute` : execute given function after given awaitable is ready
  - `bulk_then_execute` : execute given function N times after given awaitable is ready
- Executors need to minimally implement `async_execute` only
  - Missing functions are emulated

# Executors: `async_execute`

- Example implementation using futures:

```
template <typename Executor, typename F, typename ... Ts>
auto async_execute(Executor&& exec, F&& f, Ts&&... ts)
{
    hpx::promise<std::invoke_result_t<F, Ts...>> p;
    auto f = p.get_future();
    exec.sched.launch([=, p = std::move(p)]() {           // copy arguments for brevity
        p.set_value(std::invoke(f, ts...));               // assume non-void return value
    });
    return f;
}
```

# Executors: `async_execute`

- Example implementation using senders/receivers

```
template <typename Executor, typename F, typename ... Ts>
auto async_execute(Executor&& exec, F&& f, Ts&&... ts)
{
    return
        hpxexp::on(exec.sched)
        | hpxexp::then([=]() { return std::invoke(f, ts...); }));
}
```

# Executors: `bulk_async_execute`

- Example implementation agnostic to underlying execution machinery:

```
template <typename Executor, typename Shape, typename F, typename ... Ts>
auto bulk_async_execute(Executor&& exec, Shape const& shape, F&& f, Ts&&... ts)
{
    std::vector<decltype(async_execute(f, 0, ts...))> results;
    results.reserve(shape);
    for (size_t i : range(0, shape))
        results.push_back(async_execute(exec, f, i, ts...));
    return when_all(results);
}
```



# Executors: `bulk_async_execute`

- Example implementation specific to senders/receivers:

```
template <typename Executor, typename Shape, typename F, typename ... Ts>
auto bulk_async_execute(Executor&& exec, Shape const& shape, F&& f, Ts&&... ts)
{
    return
        hpxexp::on(exec.sched)
        | hpxexp::bulk(shape, [=](auto idx) { std::invoke(f, idx, ts...); });
}
```

# Parallelize Loops: Executors

Sequence of elements:

|        |   |        |   |     |        |   |
|--------|---|--------|---|-----|--------|---|
| 0      | 1 | 2      | 3 | ... | N-1    | N |
| Core 0 |   | Core 1 |   |     | Core M |   |

```
std::vector<int> d = {...};
for_each(par, d.begin(), d.end(), [](int val) {...});
```

---

```
template <typename Iterator, typename F>
auto for_each(parallel_policy policy, Iterator begin, Iterator end, F f)
{
    auto num_chunks = calculate_number_of_chunks(policy, begin, end);
    auto chunk_size = (end - begin) / num_chunks; // assume: cleanly divisible
    return wait_all(
        bulk_async_execute(
            policy.executor(), num_chunks,
            [=](size_t idx) {
                auto start_idx = chunk_size * idx;
                std::for_each(begin + start_idx, begin + start_idx + chunk_size, f); // sequential execution of chunks
            }));
}
```

# Parallelize Loops: NUM\_CHUNKS

Sequence of elements:


|        |   |        |   |     |        |   |
|--------|---|--------|---|-----|--------|---|
| 0      | 1 | 2      | 3 | ... | N-1    | N |
| Core 0 |   | Core 1 |   |     | Core M |   |

```
std::vector<int> d = {...};
hpx::for_each(with_number_of_chunks(par, NUM_CHUNKS), d.begin(), d.end(), [](int val) {...});
```

---

```
template <typename ExPolicy, typename Iterator, typename F>
auto for_each(ExPolicy&& policy, Iterator begin, Iterator end, F f)
{
    auto num_chunks = calculate_number_of_chunks(policy, begin, end);           // extract NUM_CHUNKS if given
    auto chunk_size = (end - begin) / num_chunks;                             // assume: cleanly divisible
    return bulk_async_execute(
        policy.executor(), num_chunks,
        [=](size_t idx) {
            auto start_idx = chunk_size * idx;
            hpx::for_each(begin + start_idx, begin + start_idx + chunk_size, f); // sequential execution of chunks
        });
}
```

## hpx::for\_each, memory bound speedup, Intel Skylake 20 cores (Senders & Receivers)



# Explicit Vectorization

# Vectorize Loops (explicitly)

Sequence of elements (trivial types):

|      |   |      |   |     |      |   |
|------|---|------|---|-----|------|---|
| 0    | 1 | 2    | 3 | ... | N-1  | N |
| simd |   | simd |   |     | simd |   |


```
std::vector<int> d = {...};
std::for_each(stdexp::simd, d.begin(), d.end(), [](auto val) {...});
```

---

```
template <typename Iterator, typename F>
void for_each(stdexp::simd_policy, Iterator b, Iterator e, F f)
{
    using V = stdexp::simd<Iterator::value_type>;
    for (/**/; std::distance(b, e) != 0; b += V::size()) {           // Iterator is assumed to be contiguous
        V tmp(std::addressof(*first), aligned);
        f(tmp);
        if constexpr (is_function_argument_mutable_v<F, V>)
            store(tmp, std::addressof(*first), aligned);
    }
}
```

# Vectorize Loops (explicitly)

- Provided lambda is called with a `stdexp::simd` type instance instead of a single value (see Parallelism TS V2)
- The `stdexp::simd` type has operators overloaded to make code transition seamless
- Iterator should be contiguous access
  - Otherwise bad things may happen
- Parallel algorithms load underlying sequence into vector register types before invoking loop body
- HPX implements `simd` and `par_simd` policies and their asynchronous variations



# Linear Algebra

P1673: A free function linear algebra interface based on the BLAS



# Linear Algebra

- P1673: A free function linear algebra interface based on the BLAS
  - Proposes a C++ Standard Library dense linear algebra interface

```
std::vector x_vec = { 1.0, 2.0, 3.0, ... };    // size: N

std::mdspan x(x_vec.data(), N);                // as of C++23

stdexp::linalg::scale(2.0, x);                 // sequential: x = 2.0 * x
stdexp::linalg::scale(stdexp::par, 3.0, x);    // parallel: x = 3.0 * x
```

# Linear Algebra

- Adding (optional) execution policies to all API functions
  - Allows for customization
  - Reference implementation available: <https://github.com/kokkos/stdBLAS>
    - CPU based implementation
    - Kokkos based implementation
    - HPX based implementation (under development)

# Linear Algebra

- P1673: A free function linear algebra interface based on the BLAS
  - Proposes a C++ Standard Library dense linear algebra interface

```
std::vector x_vec = { 1.0, 2.0, 3.0, ... };    // size: N

std::mdspan x(x_vec.data(), N);                // as of C++23

stdexp::linalg::scale(2.0, x);                 // sequential: x = 2.0 * x
stdexp::linalg::scale(stdexp::par, 3.0, x);    // parallel: x = 3.0 * x

stdexp::linalg::scale(hpx::par, 3.0, x);       // parallel (HPX): x = 3.0 * x
stdexp::linalg::scale(hpx::par_simd, 3.0, x);  // parallel and vectorized: x = 3.0 * x
```

# Linear Algebra: `linalg::scale` (1D)

- Exemplar 1D implementation of policy-based `linalg::scale`

```
std::vector<double> data = { 1.0, 2.0, 3.0, ... };
std::linalg::scale(par, 4.0, std::mdspan(data.data(), data.size()));
```

---

```
template <typename ExPolicy, typename Scalar, typename MdSpan>
auto scale(ExPolicy&& policy, Scalar alpha, MdSpan x)
{
    if constexpr (!supports_vectorization_v<ExPolicy> ||
                  !allow_vectorization_v<MdSpan>) {           // more conditions may apply
        // fall back to non-vectorized execution
        return hpx::for_each(to_non_simd(policy),
                             mditerator_begin(x), mditerator_end(x),
                             [&](auto& v) { v *= alpha; });
    } else {
        // possibly explicitly vectorized execution
        return hpx::for_each(policy,
                             mditerator_begin(x), mditerator_end(x),
                             [&](auto& v) { v *= alpha; });
    }
}
```

# Linear Algebra: `linalg::scale` (2D)


- Exemplar 2D implementation of policy-based `linalg::scale`

```
std::vector<double> data = { 1.0, 2.0, 3.0, ... };    // size: Nx * Ny
std::linalg::scale(par_simd, 4.0, std::mdspan(data.data(), Nx, Ny));
```

---

```
template <typename ExPolicy, typename Scalar, typename MdSpan>
auto scale(ExPolicy&& policy, Scalar alpha, MdSpan x)
{
    return hpx::for_each(to_non_simd(policy),           // allow for outer loop to be parallelized
        mditerator_begin(x), mditerator_end(x),
        [&](auto&& sub_x)
        {
            hpx::for_each(to_seq(policy),               // assume inner loop is vectorizable
                mditerator_begin(sub_x), mditerator_end(sub_x),
                [&](auto& v) { v *= alpha; });
        });
}
```

linalg::scale (1D), AMD EPYC 7352, 48 cores, 8 Vector lanes



# Conclusions

- Using execution policies for API functions that should allow for customization of execution is a good choice
  - More customization is needed, though
    - Chunking, execution environment, number of cores, etc.
  - Having means of running things asynchronously is important
    - Big hopes for senders/receivers
- Adding higher-level APIs that integrate well with senders/receivers is a must
  - Senders/receivers are fairly low level facilities with a steep learning curve
- Currently new APIs for parallel algorithms in the context of sender/receivers are being discussed
  - We believe that no new APIs are necessary



